1. Let $G = (V, E)$ be a graph. A set of edges $M \subseteq E$ is said to be a matching if no two edges in M intersect at a vertex. A matching M is *perfect* if every vertex in V is incident to some edge in M; alternatively M is perfect if $|M| = |V|/2$ (which in particular implies $|V|$ is even). See Wikipedia article for some example graphs and further background.

   The PerfectMatching problem is the following: does the given graph G have a perfect matching? This can be solved in polynomial time which is a fundamental result in combinatorial optimization with many applications in theory and practice. It turns out that the PerfectMatching problem is easier to solve in *bipartite* graphs. A graph $G = (V, E)$ is bipartite if its vertex set V can be partitioned into two sets L, R (left and right say) such that all edges are between L and R (in other words L and R are independent sets). Here is an attempted reduction from general graphs to bipartite graphs.

   Given a graph $G = (V, E)$ create a bipartite graph $H = (V \times \{1, 2\}, E_H)$ as follows. Each vertex u is made into two copies $(u, 1)$ and $(u, 2)$ with $V_1 = \{(u, 1) \mid u \in V\}$ as one side and $V_2 = \{(u, 2) \mid u \in V\}$ as the other side. Let $E_H = \{((u, 1), (v, 2)) \mid (u, v) \in E\}$. In other words we add an edge betwen $(u, 1)$ and $(v, 2)$ iff (u, v) is an edge in E. Note that $((u, 1), (u, 2))$ is not an edge in H for any $u \in V$ since there are no self-loops in G.

   Is the preceding reduction correct? To prove it is correct we need to check that H has a perfect matching if and only if G has one.

   - Prove that if G has perfect matching then H has a perfect matching.

   - Consider G to be K_3 the complete graph on 3 vertices (a triangle). Show that G has no perfect matching but H has a perfect matching.

   - Extend the previous example to obtain a graph G with an even number of vertices such that G has no perfect matching but H has.

   Thus the reduction is incorrect although one of the directions is true.

   > **Solution:**   • Suppose M is a perfect matching in G. We construct a perfect matching M' in H as follows. For each edge $(u, v) \in M$ include the edges $((u, 1), (v, 2))$ and $((v, 1), (u, 2))$ in M'. It remains to verify that M' is a perfect matching. For each vertex $u \in V$ there is exactly one edge $(u, v) \in M$ incident to u since M is perfect. This means that in H there is exactly one edge $((u, 1), (v, 2))$ incident to vertex $(u, 1)$ in M' and also exactly one edge $((v, 1), (u, 2))$ incident to $(u, 2)$. Thus M' is a perfect matching in H.
   >
   > • Suppose the three vertices of $G = K_3$ are labeled a, b, c. The edges of G are $(a, b), (b, c), (c, a)$. Then in H one can easily verify that the set of edges $\{((a, 1), (b, 2)), ((b, 1), (c, 2)), ((c, 1), (a, 2))\}$ forms a perfect matching. G does not have a perfect matching since it has an odd number of vertices.
   >
   > • Take G to be the disjoint union of two triangles, say G_1 and G_2. G has 6 vertices which is even. From the preceding part one can see that the graph H obtained from the reduction has a perfect matching since we can take the perfect matching in H_1 corresponding to G_1 and the perfect matching in H_2 corresponding to G_2 and take their union.
   >
   > ∎

2. The traveling salesman problem can be defined in two ways:

   - The Traveling Salesman Problem
     - INPUT: A weighted graph G
     - OUTPUT: Which tour $(v_1, v_2, \ldots, v_n)$ minimizes $\sum_{i=1}^{n-1} (d[v_i, v_i + 1]) + d[v_n, v_1]$
   - The Traveling Salesman *Decision* Problem
     - INPUT: A weighted graph G and an integer k
     - OUTPUT: Does there exist and TSP tour with cost $\leq k$

Suppose we are given an algorithm that can solve the traveling salesman decision problem in (say) linear time. Give an efficient algorithm to find the actual TSP tour by making a polynomial number of calls to this subroutine.

> **Solution:** We will first find the minimal cost of a TSP tour, and then find the minimal TSP tour using the minimal cost. We let $TSPD(G, k)$ denote the result of running the given algorithm for traveling salesman decision problem on the input (G, k).
>
>   **Finding the minimal cost :** Let d_{min} and d_{max} denote the minimal and maximal weight of the edges in E, respectively. Since a TSP tour would contain $|V|$ edges, we can infer that the cost of any TSP tour would be lower-bounded by $|V|d_{min}$ and upper-bounded by $|V|d_{max}$. Therefore, we can run binary search on the integers ranging from $|V|d_{min}$ to $|V|d_{max}$ to find the integer value k such that $TSPD(G, k)$ is True and $TSPD(G, k-1)$ is False. The integer k would be the minimal cost of a TSP tour.
>
>   **Finding the minimal TSP tour :** Let k be the minimal cost of a TSP tour in G. Let e be an arbitrary edge of G and let G' be a graph obtained by removing e from G. Suppose e was not included in a minimal TSP tour T of G (note that minimal TSP tour is not necessarily unique). Then $TSPD(G', k)$ would return True, since the minimal TSP tour T still remains in G'. Now suppose e was included in every minimal TSP tour of G. Then $TSPD(G', k)$ would return False, since without e we cannot formulate a minimal TSP tour with the cost k. Therefore, for each $e \in E$, we will remove e from G and run $TSPD(G, k)$. If $TSPD(G, k)$ returns True, that means there exists a minimal TSP in G that does not contain e, so we will leave e removed. If $TSPD(G, k)$ returns False, that means e is definitely a part of any minimal TSP in G, so we will add e back into G. Once we have all unnecessary edges removed, we can run a linear time search algorithm from any node to construct the tour. The following is the pseudocode of the algorithm.
>
> ```
> TSP(G(V, E)):
>     lower ← |V| * d_min
>     upper ← |V| * d_max
>     k ← BINSEARCH(lower, upper)
>     for e in E:
>         remove e from G
>         if TSPD(G, k)=False:
>             add e to G
>     s ← ARBITRARYNODE(G)
>     tour ← DFS(s)
>
>     return tour
> ```

```
BINSEARCH(lower, upper)):
    if upper − lower < 1000
        find k via brute force
    else
        mid ← (lower+upper)/2
        if TSPD(G, mid) is true
            k ← BINSEARCH(lower,mid)
        else
            k ← BINSEARCH(mid+1,upper)
    return k
```

The complexity of the binary search is dependent on the difference between the upper and lower bound of the total cost. That is, if we let $c = d_{max} - d_{min}$, the complexity of the binary search is $O(\log(c|V|))$. However, note that $c = O(2^n)$ where n is the length(number of bits) of the input, since with n bits you can only represent integers upto 2^n. Therefore, the complexity of the binary search is $O(\log(c|V|)) = O(\log c) + O(\log |V|) = O(n)$, which is linear in the input size. The rest of the algorithm loops over every edge, and TSPD is assumed to be linear ($O(V + E)$), so the runtime of the tour construction is $O(E(E + V))$. Therefore, the total runtime would be $O(n + E(E + V))$ which is polynomial in input size. For grading, the reasoning for the complexity of binary search is not required; providing $O(E(E + V))$ and reasoning about the tour construction is sufficient for full marks.

■

3. An **independent set** in a graph G is a subset S of the vertices of G, such that no two vertices in S are connected by an edge in G. Suppose you are given a magic black box that somehow answers the following decision problem *in polynomial time*:

   - INPUT: An undirected graph G and an integer k.

   - OUTPUT: TRUE if G has an independent set of size k, and FALSE otherwise.

   (a) Using this black box as a subroutine, describe algorithms that solves the following optimization problem *in polynomial time*:

   - INPUT: An undirected graph G.

   - OUTPUT: The size of the largest independent set in G.

   [Hint: You've seen this problem before.]

   > **Solution:** Suppose INDSET(V, E, k) returns TRUE if the graph (V, E) has an independent set of size k, and FALSE otherwise. Then the following algorithm returns the size of the largest independent set in G:
   >
   > > MAXINDSETSIZE(V, E):
   > >   for $k \leftarrow 1$ to V
   > >   if INDSET$(V, E, k+1) =$ FALSE
   > >   return k
   >
   > A graph with n vertices cannot have an independent set of size larger than n, so this algorithm must return a value. If G has an independent set of size k, then it also has an independent set of size $k - 1$, so the algorithm is correct.
   >
   > The algorithm clearly runs in polynomial time. Specifically, if INDSET(V, E, k) runs in $O((V + E)^c)$ time, then MAXINDSETSIZE(V, E) runs in $O((V + E)^{c+1})$ time.
   >
   > Yes, we could have used binary search instead of linear search. Whatever. ∎

(b) Using this black box as a subroutine, describe algorithms that solves the following search problem *in polynomial time*:

- INPUT: An undirected graph G.
- OUTPUT: An independent set in G of maximum size.

Solution (delete vertices): I'll use the algorithm MAXINDSETSIZE(V, E) from part (a) as a black box instead. Let $G - v$ denote the graph obtained from G by deleting vertex v, and let $G - N(v)$ denote the graph obtained from G by deleting v and all neighbors of v.

MAXINDSET(G):
  $S \leftarrow \emptyset$
  $k \leftarrow$ MAXINDSETSIZE(G)
  While G is not empty
      v is an arbitrary vertex of G
      if MAXINDSETSIZE$(G - v) = k - |S|$
          $G \leftarrow G - v$
      else
          $G \leftarrow G - N(v)$
          add v to S
  return S

Correctness of this algorithm follows inductively from the following claims:

Claim 1. *MAXINDSETSIZE$(G - v) = k$ if and only if G has an independent set of size k that excludes v.*

Proof: Every independent set in $G - v$ is also an independent set in G; it follows that MAXINDSETSIZE$(G - v) \le k$.

Suppose G has an independent set S of size k that does excludes v. Then S is also an independent set of size k in $G - v$, so MAXINDSETSIZE$(G - v)$ is at least k, and therefore equal to k.

On the other hand, suppose $G - v$ has an independent set S of size k. Then S is also a maximum independent set of G (because $|S| = k$) that excludes v. $\square$

The algorithm clearly runs in polynomial time.   ∎

ECE 374 B

Solution (add edges): I'll use the algorithm MaxIndSetSize(V, E) from part (a) as a black box instead. Let $G + uv$ denote the graph obtained from G by adding edge uv.

```
MaxIndSet(G):
    k ← MaxIndSetSize(G)
    if k = 1
        return any vertex
    for all vertices u
        for all vertices v
            if u ≠ v and uv is not an edge
                if MaxIndSetSize(G + uv) = k
                    G ← G + uv
    S ← ∅
    for all vertices v
        if deg(v) < V − 1
            add v to S
    return S
```

The algorithms adds every edge it can without changing the maximum independent set size. Let G' denote the final graph. Any independent set in G' is also an independent set in the original input graph G. Moreover, the *largest* independent set in G' is also a largest independent set in G. Thus, to prove the algorithm correct, we need to prove the following claims about the final graph G':

Claim 2. *The maximum independent set in G' is unique.*

Proof: Suppose the final graph G' has more than two maximum independent sets A and B. Pick any vertex $u \in A \setminus B$ and any other vertex $v \in A$. The set B is still an independent set in the graph $G' + uv$. Thus, when the algorithm considered edge uv, it would have added uv to the graph, contradicting the assumption that A is an independent set. □

Claim 3. *Suppose $k > 1$. The unique maximum independent set of G' contains vertex v if and only if $\deg(v) < V − 1$.*

Proof: Let S be the unique maximum independent set of G', and let v be any vertex of G. If $v \in S$, then v has degree at most $V − k < V − 1$, because v is disconnected from every other vertex in S.

On the other hand, suppose $\deg(v) < V − 1$ but $v \notin S$. Then there must be at least vertex u such that uv is not an edge in G'. Because $v \notin S$, the set S is still an independent set in $G' + uv$. Thus, when the algorithm considered edge uv, it would have added uv to the graph, and we have a contradiction. □

The algorithm clearly runs in polynomial time. ∎

To think about later:

4. Formally, a ***proper coloring*** of a graph $G = (V, E)$ is a function $c : V \to \{1, 2, \ldots, k\}$, for some integer k, such that $c(u) \neq c(v)$ for all $uv \in E$. Less formally, a valid coloring assigns each vertex of G a color, such that every edge in G has endpoints with different colors. The ***chromatic number*** of a graph is the minimum number of colors in a proper coloring of G.

   Suppose you are given a magic black box that somehow answers the following decision problem *in polynomial time*:

   - INPUT: An undirected graph G and an integer k.
   - OUTPUT: TRUE if G has a proper coloring with k colors, and FALSE otherwise.

   Using this black box as a subroutine, describe an algorithm that solves the following ***coloring problem*** in *polynomial time*:

   - INPUT: An undirected graph G.
   - OUTPUT: A valid coloring of G using the minimum possible number of colors.

   *[Hint: You can use the magic box more than once. The input to the magic box is a graph and **only** a graph, meaning **only** vertices and edges.]*

   > **Solution:** First we build an algorithm to compute the minimum number of colors in any valid coloring.
   >
   > $$\begin{array}{l} \underline{\text{CHROMATICNUMBER}(G):} \\ \quad \text{for } k \leftarrow V \text{ down to } 1 \\ \qquad \text{if COLORABLE}(G, k-1) = \text{FALSE} \\ \qquad\quad \text{return } k \end{array}$$
   >
   > Given a graph $G = (V, E)$ with n vertices $v_1, v_2, \ldots, v_n$, the following algorithm computes an array $color[1 .. n]$ describing a valid coloring of G with the minimum number of colors.

```
Coloring(G):
    k ← ChromaticNumber(G)

    ⟨⟨—— add a disjoint clique of size k ——⟩⟩
    H ← G
    for c ← 1 to k
        add vertex z_c to G
        for i ← 1 to c − 1
            add edge z_i z_c to H

    ⟨⟨—— for each vertex, try each color ——⟩⟩
    for i ← 1 to n
        for c ← 1 to k
            add edge v_i z_c to H
        for c ← 1 to k
            remove edge v_i z_c from H
            if Colorable(H, k) = True
                color[i] ← c
                break inner loop
            add edge v_i z_c from H

    return color[1 .. n]
```

In any k-coloring of H, the new vertices $z_1, \ldots, z_k$ must have k distinct colors, because every pair of those vertices is connected. We assign $color[i] \leftarrow c$ to indicate that there is a k-coloring of H in which v_i has the same color as z_c. When the algorithm terminates, $color[1 .. n]$ describes a valid k-coloring of G.

To prove that the algorithm is correct, we must prove that for all i, when the ith iteration of the outer loop ends, G has a valid k-coloring that is consistent with the partial coloring $color[1 .. i]$. Fix an integer i. The inductive hypothesis implies that when the ith iteration of the outer loop *begins*, G has a k-coloring consistent with the first $i - 1$ assigned colors. (The base case $i = 0$ is trivial.) If we connect v_i to every new vertices except z_c, then v_i must have the same color as z_c in any valid k-coloring. Thus, the call to Colorable inside the inner loop returns True if and only if H has a k-coloring in which v_i has the same color as z_c (and the previous $i - 1$ vertices are also colored). So Colorable must return True during the second inner loop, which completes the inductive proof.

This algorithm makes $O(kn) = O(n^2)$ calls to Colorable, and therefore runs in polynomial time. ∎